

Predictive Modeling of Car Accident Severity in the USA from 2016-2023 using Machine Learning Techniques

[REDACTED]

November 2023

1 Introduction

As a student, who got their driver's license only a few months ago, I am eager, yet cautious to drive. Since I learned that teen drivers have a fatal crash rate almost three times as high as drivers ages 20 and older per mile driven (CDC), I try to keep my driving to a minimum. However, avoiding driving completely is not a practical solution. Being able to drive is an integral skill of our society. Moreover, to drive safely, experience is necessary, experience that I would lack by avoiding driving entirely.

Thus, to limit my driving risk, I have decided to derive a mathematical model to evaluate and advise me on the risk of driving in various driving scenarios. Through this mathematical model, I hope to understand the risks of driving more thoroughly in order to be safer and have more peace of mind when I am on the road.

This IA will also significantly extend my mathematical understanding of how advanced statistics, linear algebra, and multivariable calculus can be used to understand and visualize trends in data, determine the most important features through dimensionality reduction, and create an accurate and applicable Machine Learning model to evaluate the driving risk of a particular situation.

1.1 Research Question

What conditions result in the most dangerous driving accidents?

2 Dataset Overview

Throughout this paper, I will be using the US Accidents (2016 - 2023) Dataset (Moosavi, 2023). Before starting any analysis, it is essential to understand what data is provided. The data columns are described in appendix A.1, appendix A.2, appendix A.3, appendix A.4, and appendix A.5. All data processing, analysis, and visualization are performed in Python appendix A.6.

3 Data Cleaning and Preprocessing

Before analyzing this dataset, it must be cleaned and processed in order to derive accurate insights and predictive models.

3.1 Reporting Sources

The data was compiled from two sources, Bing and MapQuest. These sources report severity differently, thus I had to decide which source to utilize.




Figure 1: Accident Duration by Severity




Figure 2: Impacted Distance by Severity

As can be seen by Figures 1 and 2, MapQuest has a clearly higher quartile 1, quartile 2 (median), and quartile 3 for severity 4. Bing seems to have a looser definition as the box and whisker plots are not too statistically significant from each other. For these reasons, I decided to only utilize the MapQuest data.

3.2 Non-Predictive Features

Data for columns 'ID', 'Distance(mi)', 'End Time', 'Duration', 'End Latitude', and 'End Longitude' can only be collected after the accident has already happened and hence cannot be used to predict severe accidents. Furthermore, since categorical variables like 'Country' and 'Turning Loop' only have one class, they cannot be predictive features as the probabilistic surprise and entropy for these random variables is 0 (Lesne, 2014).

3.3 Filtering Incomplete Data

Though this dataset is mostly complete, some entries lack data for certain columns. Rather than removing these columns for all entries, I decided to remove the incomplete

row entries instead as the original dataset contains over 7.7 million accidents. Even after this filtering, there is still ample data.

4 Understanding and Exploring Data

Based on Figures 1 and 2, the MapQuest accidents with severity 4 are much more serious than other severity accidents. These other levels of severity are hard to distinguish from each other. Thus, I decided to focus on severity 4 accidents (from now referred to as "severe") and group the other severity accidents (from now referred to as "non-severe") together.

4.1 Data Sample Normalization

In this processed dataset, there are roughly 2.61 million non-severe accidents and only 9000 severe accidents. This discrepancy must be normalized before conducting any further exploratory analysis. I created a Python function to return a random sample of 50000 undersampled non-severe and 50000 oversampled severe accidents.

4.2 Trends based on Time

Next, I decided to explore if there were trends based on time.

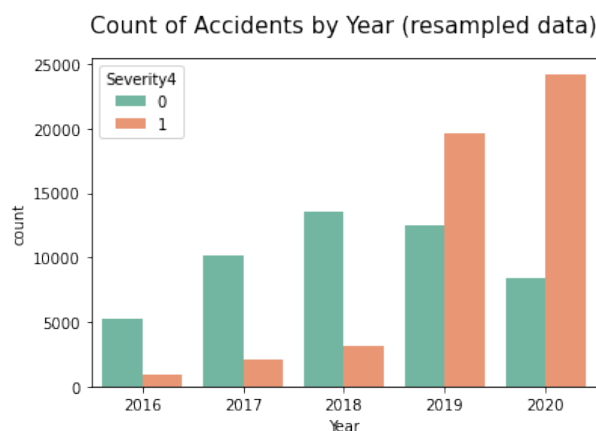


Figure 3: Count of Accidents by Year

Looking at Figure 3, it is highly improbable that severe accidents rose by 5 times from 2018 to 2019. To investigate this, I created a heatmap of these severe accidents to see how the data is actually distributed.




Figure 4: Time Heatmap of Severe Accidents

This heatmap strongly indicates that something changed after February 2019, such as the way that MapQuest defines severity or the way they collected data. Since the data after February 2019 is consistent with data in the future, dropping the data before March 2019 is the best choice for analysis and predictions.

4.3 Log Frequency Normalization

Next, I decided to analyze the accidents per hour. In Figure 5, there are two clear peaks in non-severe accidents occurring at roughly at 7-8 am and 4-5 pm, which likely correlate to maximum traffic times due to work-home commutes.




Figure 5: Count of Accidents by Hour

It is also interesting to note that during the hours (early morning or late night) where there were less non-severe accidents, there tended to be relatively greater severe accidents. This observation could be made clearer by using other data for a specific day, hour, or minute.

Since there are too many classes for a categorical variable for hours, I took the frequency of accidents during the hour. However, since certain hours could have a relatively low frequency, while others could have a relatively high frequency, the log of the frequency is taken to normalize the random continuous variable into a Gaussian distribution. Just in case the frequency is 0, 1 is added to all frequencies before taking the logarithm to prevent undefined values. This log frequency normalization (MaCurdy and Pencavel, 1986) is defined in Equation 1.

$$X = \log \left(\left(\frac{n_i}{\sum_j^k n_j} \times N_u \right) + 1 \right) \quad (1)$$

where:

- X is the resulting continuous random variable after normalization.
- n_i is the number of times a possible class of a categorical variable occurs.
- $\sum_j^k n_j$ is the total number of samples.
- N_u is the number of unique classes that the categorical variable can take.

Severe	Hour	Hour Freq
1	22	3
0	14	2
0	5	1
0	22	3
0	14	2
1	22	3

Table 1: Example Data for Hour Log Frequency Normalization

For example, using the example data in Table 1, the hour log frequency normalization for the first row can be performed using Equation 2.

$$X = \log \left(\left(\frac{n_i}{\sum_j^k n_j} \times N_u \right) + 1 \right) \quad (2)$$

$$= \log \left(\left(\frac{3}{6} \times 24 \right) + 1 \right) \quad (3)$$

$$= \log 13 \approx 2.56 \quad (4)$$

Minute Frequency by Severity (resampled data)

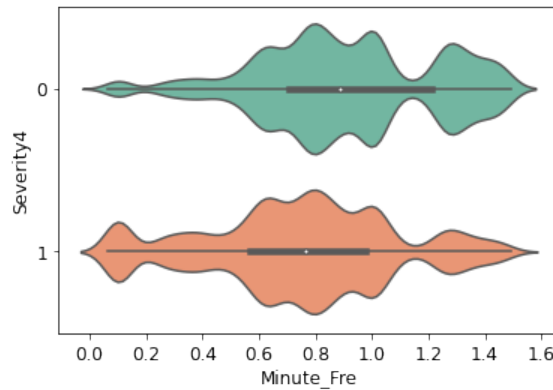


Figure 6: Accidents by Location Frequency

4.4 Location

Figure 7 shows the distribution of accidents based on latitude and longitude. The violin plots for non-severe and severe accidents seem to be very similar, though there are subtle changes in the ratio of the types of accidents.

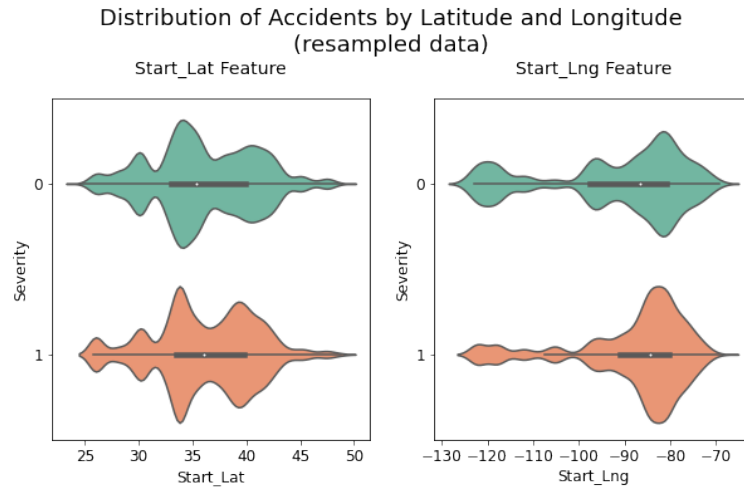


Figure 7: Accidents by Latitude and Longitude

In Figure 8, I plotted the latitude and longitude. I was fascinated to see how clearly the accidents were distributed across the US. The figure clearly shows more densely populated areas tended to have more accidents. More interesting, however, was how the severe accidents seemed to most densely be located in network-like structure. This seemed to indicate the US interstate highway system was where the greatest number of accidents seemed to occur.

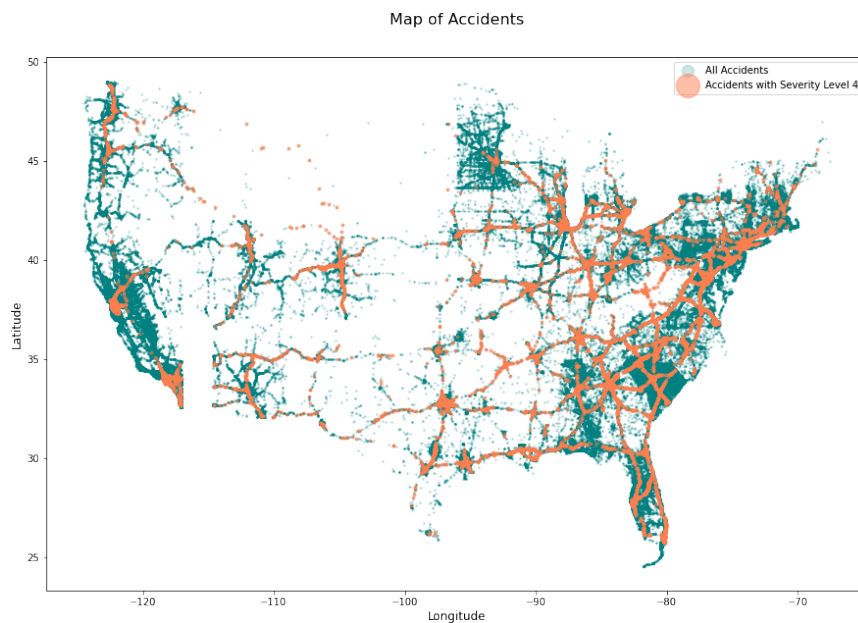


Figure 8: Map of Accidents

Using Equation 1, I used the frequency of local indicators such as 'Street', 'City', 'County', 'Zipcode', 'Airport Code', and 'State' to try to capture the clear correlation to the highway accidents in continuous random variables.

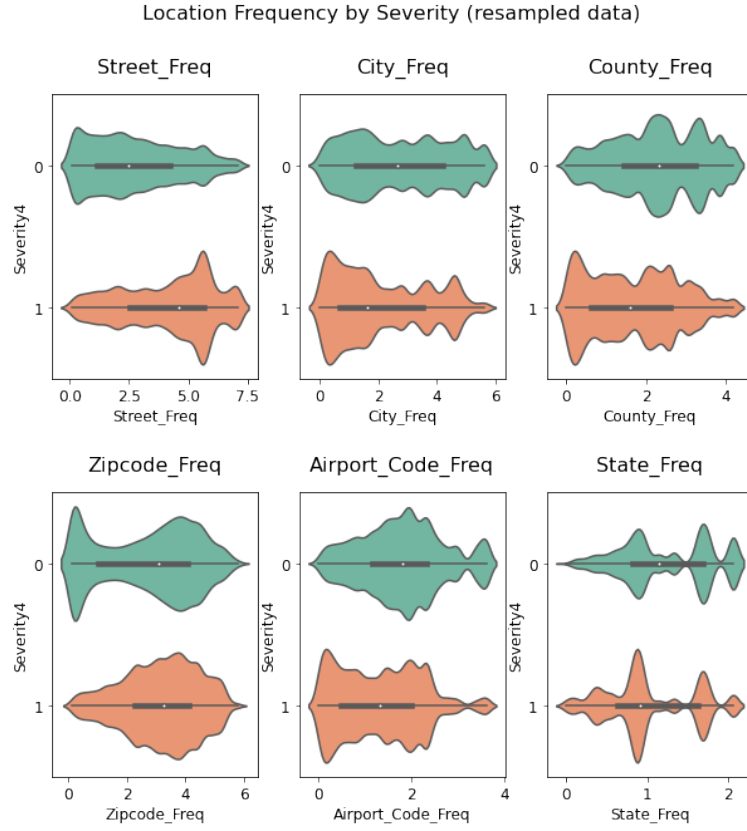


Figure 9: Accidents by Location Frequency

4.5 Weather Features

Since the weather features such as 'Pressure', 'Visibility', and 'Wind Speed' are highly skewed, they must be normalized prior to analysis. A Box Cox transformation (Osborne, 2010) can turn non-normal dependent variables into a normal Gaussian shape. By precisely controlling the parameter λ , I was able to normalize the weather features.

$$y(\lambda) = \begin{cases} \frac{y^\lambda - 1}{\lambda} & \lambda \neq 0 \\ \log y & \lambda = 0 \end{cases} \quad (5)$$

Overall, Figure 10 shows that accidents are little more likely to be serious during rain or snow while less likely on a cloudy day.

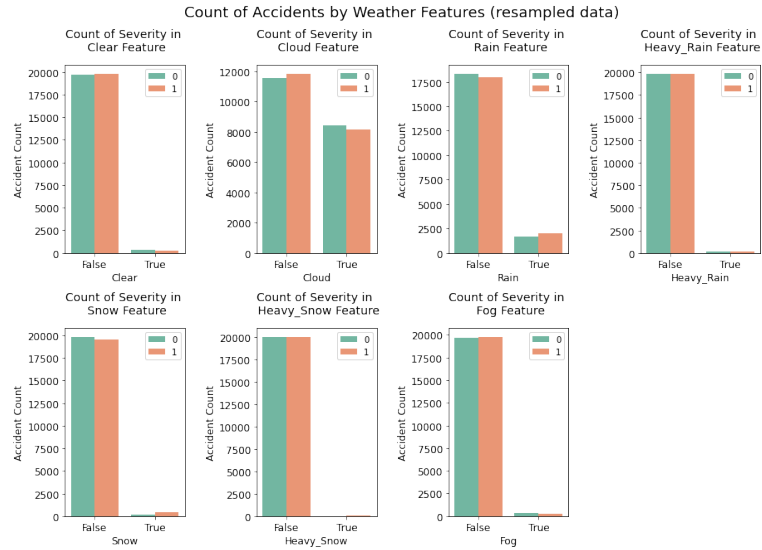


Figure 10: Accidents by Weather

5 Point of Interest Features

Figure 11 shows the number of accidents that occurred at Points of Interest (POI).

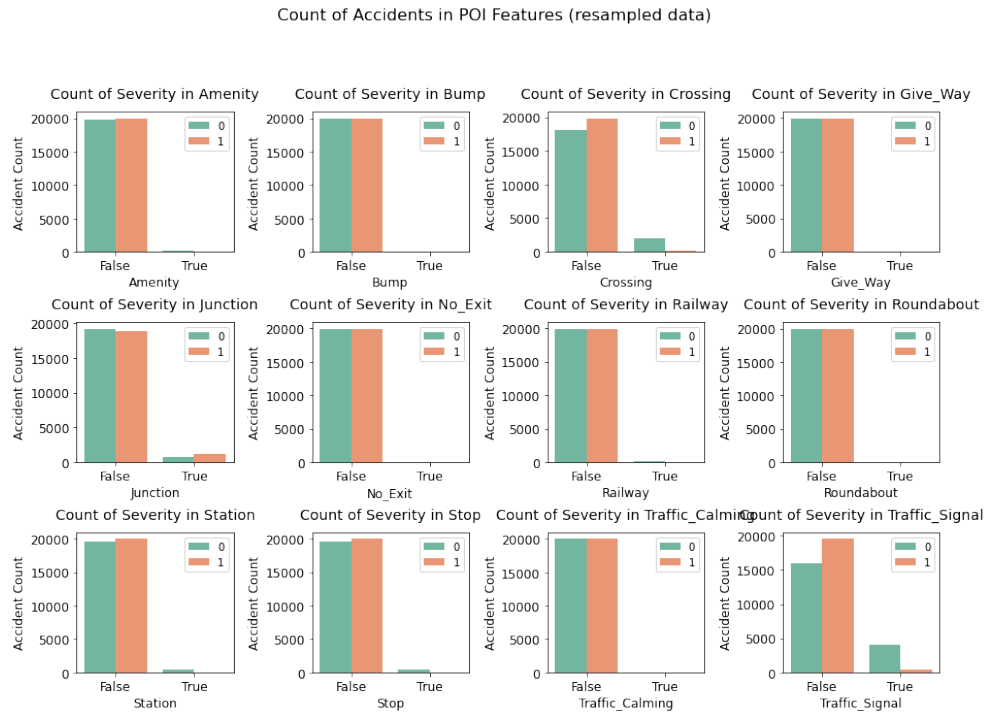


Figure 11: Accidents by Location Frequency

Accidents near traffic signal and crossing are much less likely to be serious accidents while little more likely to be serious if they are near the junction. This may occur because drivers usually slow down in front of crossing and traffic signal but junction and severity are highly related to speed. The other POI features have such little severe accidents that it

is hard to tell their relation with severity from plots.

5.1 Heatmap Correlation

After all the normalization and processing performed, the following correlation heatmap was derived.




Figure 12: Correlation Heatmap

This heatmap represents the Pearson correlation coefficient (Equation 6) for each pair of continuous random variables in the dataset.

$$r = \frac{N \sum XY - (\sum X \sum Y)}{\sqrt{[N \sum x^2 - (\sum x)^2][N \sum y^2 - (\sum y)^2]}} \quad (6)$$

The Correlation Heatmap is promising as it shows both mildly strong positive and negative correlations of the various data columns to the severity, while being minimally correlated to each other.

This should in theory should enable enable a predictive model to optimize weights of select factors in order to derive a strong fit.

6 Predictive Modeling

Using the processed data from Section 4, I finally could start on designing a predictive model.

6.1 Principal Component Analysis

Even though I had majorly reduced the dimensionality of the data from 48 columns to only 12 columns through the processing steps, this was still too much data to pass into a predictive model. As a result, I decided to use Principal Component Analysis (PCA), which reduces data dimensionality while retaining most of the relevant information.

PCA involves:

1. Centering the data about the origin.

- $X'_i = X_i - \bar{X}$

2. Calculating the covariance matrix with n variables.

- $Cov(X, Y) = \frac{1}{n-1} \sum_{i=1}^n (X'_i - \bar{X}) (Y'_i - \bar{Y})$

3. Compute the eigenvalues λ and eigenvectors v of the covariance matrix C .

- $Cv = \lambda v$

4. Selecting the top k principal component eigenvalues and corresponding eigenvectors.

5. Projecting an original point (X) onto the new feature space along the i th principle component v_i .

- $X' = Xv_i$

6. The principal components are linear combinations of the original variables. a_{ij} represents the components of the i th eigenvector. This represents the general form for the principal component i .

- $PC_i = a_{i1}X'_1 + a_{i2}X'_2 + \dots + a_{in}X'_n$

After performing PCA, the explained variance ratio can be derived using the Equation 7. Using PC1, PC2, ..., and PC8, over 80% of the variance in the accident severity is captured. This is described in the Scree Plot in Figure 13.

$$\text{Explained Variance} = \frac{\lambda_i}{\lambda_1 + \lambda_2 + \dots + \lambda_n} \quad (7)$$




Figure 13: Scree Plot

6.2 Binary Logistic Regression

In order to prevent overfitting and address the bias-variance tradeoff (Belkin et al., 2019), I split the data into 80% training and 20% testing. Using the PCs derived in the previous step, I decided to try using a binary logistic regression.

A binary logistic regression is defined by Equation 8.

$$y = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}} \quad (8)$$

To tune the parameters β_0 and β_1 , a maximum likelihood interpretation is utilized. Equation 9 describes the calculation for joint likelihood $L(\beta_0, \beta)$.

$$L(\beta_0, \beta) = \prod_{i=1}^n (p(x_i))^{y_i} \times (1 - p(x_i))^{1-y_i} \quad (9)$$

This can further be simplified by taking the log-likelihood $l(\beta_0, \beta)$ described in Equation 10

$$l(\beta_0, \beta) = \sum_{i=1}^n y_i \log(p(x_i)) + (1 - y_i) \log(1 - p(x_i)) \quad (10)$$

After taking the partial derivative of the log-likelihood $l(\beta_0, \beta)$ with respect to β_0 and β parameters individually and simplifying, Equation 11 describes the gradient of the β_0 and β_1 . By numerically stepping the β_0 and β_1 using the gradient from Equation 11, the likelihood that the data was derived from the fitted logistic regression is maximized (Menard, 2002).

$$\frac{\partial l}{\partial \beta_j} = - \sum_{i=1}^n (y_i - p(x_i; \beta_0, \beta)) x_{ij} \quad (11)$$

6.2.1 Results




Figure 14: Logistic Regression Confusion Matrix

Using Figure 14 and Equations 13, 16, and 19, I calculated the sensitivity, specificity, and accuracy of my model.

$$\text{Sensitivity} = \frac{TP}{TP + FN} \quad (12)$$

$$= \frac{7590}{7590 + 2445} \quad (13)$$

$$\approx 75.6\% \quad (14)$$

$$\text{Specificity} = \frac{TN}{TN + FP} \quad (15)$$

$$= \frac{6829}{6829 + 3136} \quad (16)$$

$$\approx 68.5\% \quad (17)$$

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + FN + TN} \quad (18)$$

$$= \frac{7590 + 6829}{7590 + 3136 + 2445 + 6829} \quad (19)$$

$$\approx 72.1\% \quad (20)$$

where:

- TP is the number of True Positives.
- TN is the number of True Negatives.

- FP is the number of False Positives.
- FN is the number of False Negatives.

As can be seen in Equations 13, 16, and 19, the model has a relatively low sensitivity, specificity, and accuracy. I believe that a stronger predictive model can be created.

6.3 Random Forest Classifier

Given relatively weak results from Section 6.2, I decided to try the more advanced Random Forest Classifier using GridSearchCV. Random Forest classifier is a machine learning model that combines the predictions of multiple decision trees to make more accurate and robust classifications. It works by creating a forest of decision trees, each trained on a random subset of the data and using a random subset of the features. The individual tree predictions are then aggregated to make the final classification decision, often through a majority vote or weighted average. Random Forests are known for their ability to handle high-dimensional data, reduce overfitting, and provide feature importance rankings, making them a popular choice for a wide range of classification tasks, from image recognition to financial risk assessment.

6.3.1 Results




Figure 15: Random Forest Confusion Matrix

$$\text{Sensitivity} = \frac{TP}{TP + FN} \quad (21)$$

$$= \frac{9859}{9859 + 176} \quad (22)$$

$$\approx 98.2\% \quad (23)$$

$$\text{Specificity} = \frac{TN}{TN + FP} \quad (24)$$

$$= \frac{8794}{8794 + 1171} \quad (25)$$

$$\approx 88.2\% \quad (26)$$

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + FN + TN} \quad (27)$$

$$= \frac{9859 + 8794}{9859 + 1171 + 176 + 8794} \quad (28)$$

$$\approx 93.3\% \quad (29)$$

Figure 15 and Equations 22, 25, and 28 clearly show that the Random Forest Model is more sensitive, specific, and accurate than the Logistic Model in Section 6.2. Since the sensitivity (true positive rate) is 98% and the specificity (true negative rate) is 88%, the Random Forest Classifier seems to miscategorize more non-severe accidents as severe (Zhu et al., 2010). This is acceptable for the practical application of this model as overestimating the severity of accidents would bring greater public awareness of dangerous driving conditions.

7 Conclusion

The most accurate predictive model that I created was the Random Forest Classifier with a 93%. Figure 16 shows the most significant feature weights that the model iteratively learned.

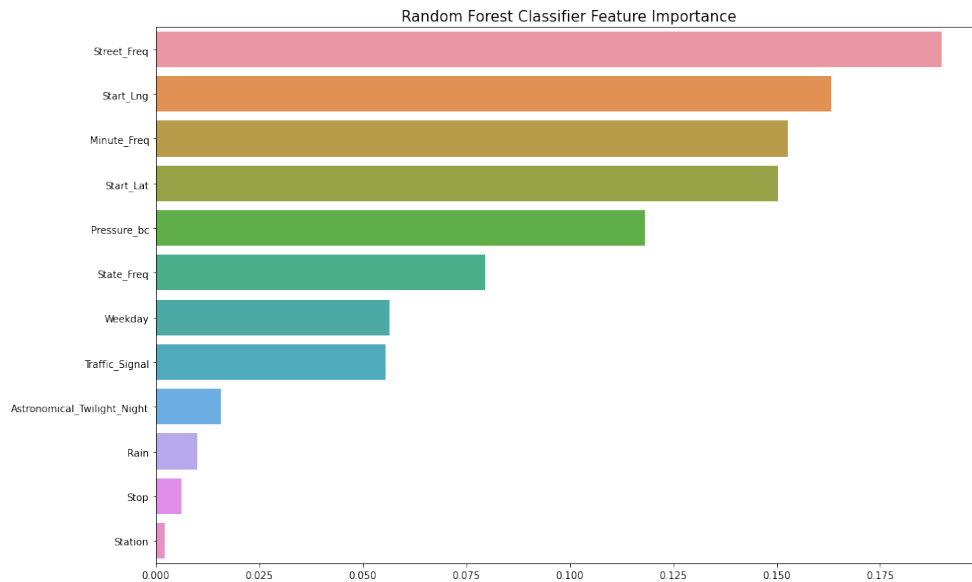


Figure 16: Random Forest Classifier Feature Importance

The 5 most important features it determined were:

1. Street Frequency
2. Start Longitude
3. Minute Frequency
4. Start Latitude
5. Pressure (Box-Cox)

I am not surprised that Street Frequency, Start Longitude, and Start Latitude were in the top 5. As seen in Section 4, most severe accidents seemed to occur on interstate highways in densely populated regions (based on longitude and latitude). Along with the temporal input that the Minute Frequency feature provided, the pattern clearly indicates that severe accidents are likely to occur in the same location and time as severe accidents that have occurred in the past.

Unlike the other factors, I was surprised by the Pressure (Box-Cox) feature being so important. Since a strong negative correlation was found for this factor, I wonder if pressure plays an indirect role in causing severe accidents. This is unlikely to be due to the weather as none of the weather features like Rain seemed to have a strong correlation with severity.

I was also surprised to see factors like Astronomical Twilight Night and Traffic Signal to only be somewhat important in predicting severe accidents as these are often cited as common dangerous driving conditions. As a further exploration, I would like to investigate these anomalies.

References

- Belkin, M., Hsu, D., Ma, S., and Mandal, S. (2019). Reconciling modern machine-learning practice and the classical bias–variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32):15849–15854.
- Lesne, A. (2014). Shannon entropy: a rigorous notion at the crossroads between probability, information theory, dynamical systems and statistical physics. *Mathematical Structures in Computer Science*, 24(3):e240311.
- MaCurdy, T. E. and Pencavel, J. H. (1986). Testing between competing models of wage and employment determination in unionized markets. *Journal of Political Economy*, 94(3):S3–S39.
- Menard, S. (2002). *Applied logistic regression analysis*. Sage.
- Moosavi, S. (2023). Us accidents (2016 - 2023).
- Osborne, J. (2010). Improving your data transformations: Applying the box-cox transformation. *Practical Assessment, Research, and Evaluation*, 15(1):12.
- Zhu, W., Zeng, N., Wang, N., et al. (2010). Sensitivity, specificity, accuracy, associated confidence interval and roc analysis with practical sas implementations. *NESUG proceedings: health care and life sciences, Baltimore, Maryland*, 19:67.

A Appendix

A.1 Traffic Attributes

Column	Description
ID	This is a unique identifier of the accident record.
Source	Indicates source of the accident report.
Severity	Shows the severity of the accident, a number between 1 and 4.
Start Time	Shows start time of the accident in local time zone.
End Time	Shows end time of the accident in local time zone.
Start Latitude	Shows latitude in GPS coordinates of the start point.
Start Longitude:	Shows longitude in GPS coordinate of the start point.
End Latitude	Shows latitude in GPS coordinate of the end point.
End Longitude	Shows longitude in GPS coordinate of the end point.
Distance(mi)	The length of the road extent affected by the accident.
Description	Shows natural language description of the accident.

A.2 Address Attributes

Column	Description
Number	Shows the street number in address field.
Street	Shows the street name in address field.
City	Shows the city in address field.
County	Shows the county in address field.
State	Shows the state in address field.
Zip Code	Shows the zipcode in address field.
Country	Shows the country in address field.
Timezone	Shows timezone based on the location of the accident.

A.3 Weather Attributes

Column	Description
Airport Code	Denotes the closest airport-based weather station.
Weather Timestamp	Shows the time-stamp of weather observation.
Temperature	Shows the temperature (in Fahrenheit).
Wind Chill	Shows the wind chill (in Fahrenheit).
Humidity	Shows the humidity (in percentage).
Pressure	Shows the air pressure (in inches).
Visibility	Shows visibility (in miles).
Wind Direction	Shows wind direction.
Wind Speed	Shows wind speed (in miles per hour).

Precipitation	Shows precipitation amount in inches, if there is any.
Weather Condition	Shows the weather condition.

A.4 Point-Of-Interest Attributes

Column	Description
Amenity	Indicates presence of amenity in a nearby location.
Bump	Indicates presence of speed bump or hump in a nearby location.
Crossing	Indicates presence of crossing in a nearby location.
Give Way	Indicates presence of give way sign in a nearby location.
Junction	Indicates presence of junction in a nearby location.
No Exit	Indicates presence of no exit sign in a nearby location.
Railway	Indicates presence of railway in a nearby location.
Roundabout	Indicates presence of roundabout in a nearby location.
Station	Indicates presence of station (bus, train, etc.) in a nearby location.
Stop	Indicates presence of stop sign in a nearby location.
Traffic Calming	Indicates presence of traffic calming means in a nearby location.
Traffic Signal	Indicates presence of traffic signal in a nearby location.
Turning Loop	Indicates presence of turning loop in a nearby location.

A.5 Period-of-Day Attributes

Column	Description
Sunrise Sunset	Shows the period of day based on sunrise/sunset.
Civil Twilight	Shows the period of day based on civil twilight.
Nautical Twilight	Shows the period of day based on nautical twilight.
Astronomical Twilight	Shows the period of day based on astronomical twilight.

A.6 Python Code for Data Processing, Analysis, and Visualization

```

1 import numpy as np
2 import pandas as pd
3 import json
4 import matplotlib
5 import matplotlib.pyplot as plt
6 from matplotlib import cm
7 from datetime import datetime
8 import glob
9 import seaborn as sns
10 import re
11 import os

```

```

12 import io
13 from scipy.stats import boxcox
14
15 df =
    ↪ pd.read_csv('../input/us-accidents/US_Accidents_March23.csv')
16 print("The shape of data is:", (df.shape))
17 print(df.head(3))
18
19 df_source =
    ↪ df.groupby(['Severity', 'Source']).size().reset_index().pivot(\
20     columns='Severity', index='Source', values=0)
21 df_source.plot(kind='bar', stacked=True, title='Severity Count by
    ↪ Sources')
22
23 # fix datetime type
24 df['Start_Time'] = pd.to_datetime(df['Start_Time'])
25 df['End_Time'] = pd.to_datetime(df['End_Time'])
26 df['Weather_Timestamp'] = pd.to_datetime(df['Weather_Timestamp'])
27
28 # calculate duration as the difference between end time and start
    ↪ time in minute
29 df['Duration'] = df.End_Time - df.Start_Time
30 df['Duration'] = df['Duration'].apply(lambda
    ↪ x:round(x.total_seconds() / 60) )
31 print("The overall mean duration is: ",
    ↪ (round(df['Duration'].mean(), 3)), 'min')
32
33 fig, axs = plt.subplots(ncols=2, figsize=(10, 4))
34 sns.boxplot(x="Severity", y="Duration",
35     data=df.loc[(df['Source']=="Source2") &
    ↪ (df['Duration']<400),], palette="Set2",
    ↪ ax=axs[0])
36 axs[0].set_title('MapQuest')
37 fig.suptitle('Accidents Duration by Severity', fontsize=16)
38 sns.boxplot(x="Severity", y="Duration",
39     data=df.loc[(df['Source']=="Source1") &
    ↪ (df['Duration']<400),], palette="Set2",
    ↪ ax=axs[1])
40 axs[1].set_title('Bing')
41 plt.show()
42

```

```

43 fig, axs = plt.subplots(ncols=2, figsize=(10, 4))
44 sns.boxplot(x="Severity", y="Distance(mi)",
45             data=df.loc[(df['Source']=="Source2") &
46                         ↪ (df['Distance(mi)']<10),], palette="Set2",
47                         ↪ ax=axs[0])
48 axs[0].set_title('MapQuest')
49 fig.suptitle('Impacted Distance by Severity', fontsize=16)
50 sns.boxplot(x="Severity", y="Distance(mi)",
51             data=df.loc[(df['Source']=="Source1") &
52                         ↪ (df['Distance(mi)']<10),], palette="Set2",
53                         ↪ ax=axs[1])
54 axs[1].set_title('Bing')
55 plt.show()
56
57 df = df.loc[df['Source']=="Source2",]
58 df = df.drop(['Source'], axis=1)
59 print("The shape of data is:", (df.shape))
60
61 df = df.drop(['ID', 'Description', 'Distance(mi)', 'End_Time',
62             ↪ 'Duration',
63             ↪ 'End_Lat', 'End_Lng'], axis=1)
64
65 cat_names = ['Country', 'Timezone', 'Amenity', 'Bump',
66             ↪ 'Crossing',
67             ↪ 'Give_Way', 'Junction', 'No_Exit', 'Railway',
68             ↪ 'Roundabout', 'Station',
69             ↪ 'Stop', 'Traffic_Calming', 'Traffic_Signal',
70             ↪ 'Turning_Loop', 'Sunrise_Sunset',
71             ↪ 'Civil_Twilight', 'Nautical_Twilight',
72             ↪ 'Astronomical_Twilight']
73 print("Unique count of categorical features:")
74 for i in cat_names:
75     print(i, df[i].unique().size)
76
77 df = df.drop(['Country', 'Turning_Loop'], axis=1)
78
79 print("Wind Direction: ", df['Wind_Direction'].unique())
80
81 df.loc[df['Wind_Direction']=='Calm', 'Wind_Direction'] = 'CALM'
82 df.loc[(df['Wind_Direction']=='West') | (df['Wind_Direction']=='WSW') | (df['Wind_Direction']=='W'), 'Wind_Direction'] = 'W'

```



```

93
94 # Assign NA to created weather features where 'Weather_Condition'
   ↳ is null.
95 weather =
   ↳ ['Clear', 'Cloud', 'Rain', 'Heavy_Rain', 'Snow', 'Heavy_Snow', 'Fog']
96 for i in weather:
97     df.loc[df['Weather_Condition'].isnull(), i] =
   ↳ df.loc[df['Weather_Condition'].isnull(), 'Weather_Condition']
98     df[i] = df[i].astype('bool')
99
100 df.loc[:, ['Weather_Condition'] + weather]
101
102 df = df.drop(['Weather_Condition'], axis=1)
103
104 # average difference between weather time and start time
105 print("Mean difference between 'Start_Time' and
   ↳ 'Weather_Timestamp': ",
106 (df.Weather_Timestamp - df.Start_Time).mean())
107
108 df = df.drop(["Weather_Timestamp"], axis=1)
109
110 df['Year'] = df['Start_Time'].dt.year
111
112 nmonth = df['Start_Time'].dt.month
113 df['Month'] = nmonth
114
115 df['Weekday'] = df['Start_Time'].dt.weekday
116
117 days_each_month =
   ↳ np.cumsum(np.array([0, 31, 28, 31, 30, 31, 30, 31, 31, 30, 31, 30, 31]))
118 nday = [days_each_month[arg-1] for arg in nmonth.values]
119 nday = nday + df["Start_Time"].dt.day.values
120 df['Day'] = nday
121
122 df['Hour'] = df['Start_Time'].dt.hour
123
124 df['Minute'] = df['Hour'] * 60.0 + df["Start_Time"].dt.minute
125
126 df.loc[:, 4, ['Start_Time', 'Year', 'Month', 'Weekday', 'Day',
   ↳ 'Hour', 'Minute']]
127

```

```

128 missing = pd.DataFrame(df.isnull().sum()).reset_index()
129 missing.columns = ['Feature', 'Missing_Percent(%)']
130 missing['Missing_Percent(%)'] =
    ↳ missing['Missing_Percent(%)'].apply(lambda x: x / df.shape[0]
    ↳ * 100)
131 missing.loc[missing['Missing_Percent(%)']>0,:]
132
133 df = df.drop(['Wind_Chill(F)'], axis=1)
134
135 df['Precipitation_NA'] = 0
136 df.loc[df['Precipitation(in)'].isnull(), 'Precipitation_NA'] = 1
137 df['Precipitation(in)'] =
    ↳ df['Precipitation(in)'].fillna(df['Precipitation(in)'].median())
138 df.loc[:5, ['Precipitation(in)', 'Precipitation_NA']]
139
140 df = df.dropna(subset=['City', 'Zipcode', 'Airport_Code',
    ↳ 'Sunrise_Sunset', 'Civil_Twilight', 'Nautical_Twilight'])
141
142
143 # group data by 'Airport_Code' and 'Start_Month' then fill NAs
    ↳ with median value
144 Weather_data=['Temperature(F)', 'Humidity(%)', 'Pressure(in)', 'Visibility(mi)',
145 print("The number of remaining missing values: ")
146 for i in Weather_data:
147     df[i] = df.groupby(['Airport_Code', 'Month'])[i].apply(lambda x:
    ↳ x.fillna(x.median()))
148     print(i + " : " + df[i].isnull().sum().astype(str))
149
150 df = df.dropna(subset=Weather_data)
151
152 # group data by 'Airport_Code' and 'Start_Month' then fill NAs
    ↳ with majority value
153 from collections import Counter
154 weather_cat = ['Wind_Direction'] + weather
155 print("Count of missing values that will be dropped: ")
156 for i in weather_cat:
157     df[i] = df.groupby(['Airport_Code', 'Month'])[i].apply(lambda x:
    ↳ x.fillna(Counter(x).most_common()[0][0]) if
    ↳ all(x.isnull())==False else x)
158     print(i + " : " + df[i].isnull().sum().astype(str))
159

```



```

160 # drop na
161 df = df.dropna(subset=weather_cat)
162
163 df['Severity4'] = 0
164 df.loc[df['Severity'] == 4, 'Severity4'] = 1
165 df = df.drop(['Severity'], axis = 1)
166 df.Severity4.value_counts()
167
168 def resample(dat, col, n):
169     return pd.concat([dat[dat[col]==1].sample(n, replace = True),
170                      dat[dat[col]==0].sample(n)], axis=0)
171
172 df_b1 = resample(df, 'Severity4', 50000)
173 print('resampled data:', df_b1.Severity4.value_counts())
174
175 df_b1.Year = df_b1.Year.astype(str)
176 sns.countplot(x='Year', hue='Severity4', data=df_b1
177               ↪ ,palette="Set2")
178 plt.title('Count of Accidents by Year (resampled data)', size=15,
179           ↪ y=1.05)
180 plt.show()
181
182 # create a dataframe used to plot heatmap
183 df_date = df.loc[:, ['Start_Time', 'Severity4']] # create a
184           ↪ new dataframe only containing time and severity
185 df_date['date'] = df_date['Start_Time'].dt.normalize() # keep
186           ↪ only the date part of start time
187 df_date = df_date.drop(['Start_Time'], axis = 1)
188 df_date = df_date.groupby('date').sum() # sum the
189           ↪ number of accidents with severity level 4 by date
190 df_date = df_date.reset_index().drop_duplicates()
191
192 # join the dataframe with full range of date from 2016 to 2020
193 full_date =
194     ↪ pd.DataFrame(pd.date_range(start="2016-01-02",end="2020-12-31"))
195 df_date = full_date.merge(df_date, how = 'left', left_on = 0,
196                          ↪ right_on = 'date')
197 df_date['date'] = df_date.iloc[:,0]
198 df_date = df_date.fillna(0)
199 df_date = df_date.iloc[:,1:].set_index('date')

```

```

194 # group by date
195 groups = df_date['Severity4'].groupby(pd.Grouper(freq='A'))
196 years = pd.DataFrame()
197 for name, group in groups:
198     if name.year != 2020:
199         years[name.year] = np.append(group.values, 0)
200     else:
201         years[name.year] = group.values
202
203
204 # plot
205 years = years.T
206 plt.matshow(years, interpolation=None, aspect='auto')
207 plt.title('Time Heatmap of Accident with Severity Level 4 (raw
↳ data)', y=1.2, fontsize=15)
208 plt.show()
209
210 df = df.loc[df['Start_Time'] > "2019-03-10",:]
211 df = df.drop(['Year', 'Start_Time'], axis=1)
212 df['Severity4'].value_counts()
213
214 df_bl = resample(df, 'Severity4', 20000)
215
216 plt.figure(figsize=(10,5))
217 sns.countplot(x='Month', hue='Severity4', data=df_bl
↳ ,palette="Set2")
218 plt.title('Count of Accidents by Month (resampled data)',
↳ size=15, y=1.05)
219 plt.show()
220
221 plt.figure(figsize=(10,5))
222 sns.countplot(x='Weekday', hue='Severity4', data=df_bl
↳ ,palette="Set2")
223 plt.title('Count of Accidents by Weedday (resampled data)',
↳ size=15, y=1.05)
224 plt.show()
225
226 period_features =
↳ ['Sunrise_Sunset', 'Civil_Twilight', 'Nautical_Twilight', 'Astronomical_Twili
227 fig, axs = plt.subplots(ncols=1, nrows=4, figsize=(13, 5))
228

```

```

229 plt.subplots_adjust(wspace = 0.5)
230 for i, feature in enumerate(period_features, 1):
231     plt.subplot(1, 4, i)
232     sns.countplot(x=feature, hue='Severity4', data=df_bl
        ↪ ,palette="Set2")
233
234     plt.xlabel('{}'.format(feature), size=12, labelpad=3)
235     plt.ylabel('Accident Count', size=12, labelpad=3)
236     plt.tick_params(axis='x', labelsiz=12)
237     plt.tick_params(axis='y', labelsiz=12)
238
239     plt.legend(['0', '1'], loc='upper right', prop={'size': 10})
240     plt.title('Count of Severity in\n{} Feature'.format(feature),
        ↪ size=13, y=1.05)
241 fig.suptitle('Count of Accidents by Period-of-Day (resampled
        ↪ data)',y=1.08, fontsize=16)
242 plt.show()
243
244 plt.figure(figsize=(15,5))
245 sns.countplot(x='Hour', hue='Severity4', data=df_bl
        ↪ ,palette="Set2")
246 plt.title('Count of Accidents by Hour (resampled data)', size=15,
        ↪ y=1.05)
247 plt.show()
248
249 # frequency encoding and log-transform
250 df['Minute_Freq'] =
        ↪ df.groupby(['Minute'])['Minute'].transform('count')
251 df['Minute_Freq'] = df['Minute_Freq']/df.shape[0]*24*60
252 df['Minute_Freq'] = df['Minute_Freq'].apply(lambda x:
        ↪ np.log(x+1))
253
254 # resampling
255 df_bl = resample(df, 'Severity4', 20000)
256
257 # plot
258 df_bl['Severity4'] = df_bl['Severity4'].astype('category')
259 sns.violinplot(x='Minute_Freq', y="Severity4", data=df_bl,
        ↪ palette="Set2")
260 plt.xlabel('Minute_Fre', size=12, labelpad=3)
261 plt.ylabel('Severity4', size=12, labelpad=3)

```

```

262 plt.tick_params(axis='x', labelsiz=12)
263 plt.tick_params(axis='y', labelsiz=12)
264 plt.title('Minute Frequency by Severity (resampled data)',
    ↪ size=16, y=1.05)
265 plt.show()
266
267 plt.figure(figsize=(6,5))
268 chart = sns.countplot(x='Timezone', hue='Severity4', data=df_bl
    ↪ ,palette="Set2")
269 plt.title("Count of Accidents by Timezone (resampled data)",
    ↪ size=15, y=1.05)
270 plt.show()
271
272 plt.figure(figsize=(25,5))
273 chart = sns.countplot(x='State', hue='Severity4',
274                      data=df_bl ,palette="Set2",
    ↪ order=df_bl['State'].value_counts().index)
275 plt.title("Count of Accidents in State\nordered by accidents'
    ↪ count (resampled data)", size=15, y=1.05)
276 plt.show()
277
278 df_bl['Severity4'] = df_bl['Severity4'].astype('category')
279 num_features = ['Start_Lat', 'Start_Lng']
280 fig, axs = plt.subplots(ncols=1, nro=2, figsize=(10, 5))
281 plt.subplots_adjust(hspace=0.4, wspace = 0.2)
282 for i, feature in enumerate(num_features, 1):
283     plt.subplot(1, 2, i)
284     sns.violinplot(x=feature, y="Severity4", data=df_bl,
    ↪ palette="Set2")
285
286     plt.xlabel('{}'.format(feature), size=12, labelpad=3)
287     plt.ylabel('Severity', size=12, labelpad=3)
288     plt.tick_params(axis='x', labelsiz=12)
289     plt.tick_params(axis='y', labelsiz=12)
290
291     plt.title('{} Feature'.format(feature), size=14, y=1.05)
292 fig.suptitle('Distribution of Accidents by Latitude and
    ↪ Longitude\n(resampled data)', fontsize=18, y=1.08)
293 plt.show()
294
295 df_4 = df[df['Severity4']==1]

```

```

296
297 plt.figure(figsize=(15,10))
298
299 plt.plot( 'Start_Lng', 'Start_Lat', data=df, linestyle='',
    ↪ marker='o', markersize=1.5, color="teal", alpha=0.2,
    ↪ label='All Accidents')
300 plt.plot( 'Start_Lng', 'Start_Lat', data=df_4, linestyle='',
    ↪ marker='o', markersize=3, color="coral", alpha=0.5,
    ↪ label='Accidents with Severity Level 4')
301 plt.legend(markerscale=8)
302 plt.xlabel('Longitude', size=12, labelpad=3)
303 plt.ylabel('Latitude', size=12, labelpad=3)
304 plt.title('Map of Accidents', size=16, y=1.05)
305 plt.show()
306
307 fre_list = ['Street', 'City', 'County', 'Zipcode',
    ↪ 'Airport_Code', 'State']
308 for i in fre_list:
309     newname = i + '_Freq'
310     df[newname] = df.groupby([i])[i].transform('count')
311     df[newname] = df[newname]/df.shape[0]*df[i].unique().size
312     df[newname] = df[newname].apply(lambda x: np.log(x+1))
313
314 # resample again
315 df_bl = resample(df, 'Severity4', 50000)
316
317 df_bl['Severity4'] = df_bl['Severity4'].astype('category')
318 fig, axs = plt.subplots(ncols=2, nrows=3, figsize=(10, 10))
319 plt.subplots_adjust(hspace=0.4, wspace = 0.2)
320 fig.suptitle('Location Frequency by Severity (resampled data)',
    ↪ fontsize=16)
321 for i, feature in enumerate(fre_list, 1):
322     feature = feature + '_Freq'
323     plt.subplot(2, 3, i)
324     sns.violinplot(x=feature, y="Severity4", data=df_bl,
    ↪ palette="Set2")
325
326     plt.xlabel('{}'.format(feature), size=12, labelpad=3)
327     plt.ylabel('Severity4', size=12, labelpad=3)
328     plt.tick_params(axis='x', labelsiz=12)
329     plt.tick_params(axis='y', labelsiz=12)

```

```

330
331     plt.title('{}'.format(feature), size=16, y=1.05)
332 plt.show()
333
334 df = df.drop(fre_list, axis = 1)
335
336 df['Pressure_bc'] = boxcox(df['Pressure(in)'].apply(lambda x:
    ↪ x+1), lmbda=0.3)
337 df['Visibility_bc'] = boxcox(df['Visibility(mi)'].apply(lambda x:
    ↪ x+1), lmbda = 0.1)
338 df['Wind_Speed_bc'] = boxcox(df['Wind_Speed(mph)'].apply(lambda x:
    ↪ x+1), lmbda=-0.2)
339 df = df.drop(['Pressure(in)', 'Visibility(mi)', 'Wind_Speed(mph)'],
    ↪ axis=1)
340
341 # resample again
342 df_bl = resample(df, 'Severity4', 50000)
343
344 df_bl['Severity4'] = df_bl['Severity4'].astype('category')
345 num_features = ['Temperature(F)', 'Humidity(%)', 'Pressure_bc',
    ↪ 'Visibility_bc', 'Wind_Speed_bc']
346 fig, axs = plt.subplots(ncols=2, nrows=3, figsize=(15, 10))
347 plt.subplots_adjust(hspace=0.4, wspace = 0.2)
348 for i, feature in enumerate(num_features, 1):
349     plt.subplot(2, 3, i)
350     sns.violinplot(x=feature, y="Severity4", data=df_bl,
    ↪ palette="Set2")
351
352     plt.xlabel('{}'.format(feature), size=12, labelpad=3)
353     plt.ylabel('Severity', size=12, labelpad=3)
354     plt.tick_params(axis='x', labelsize=12)
355     plt.tick_params(axis='y', labelsize=12)
356
357     plt.title('{} Feature by Severity'.format(feature), size=14,
    ↪ y=1.05)
358 fig.suptitle('Density of Accidents by Weather Features (resampled
    ↪ data)', fontsize=18)
359 plt.show()
360
361 fig, axs = plt.subplots(ncols=2, nrows=4, figsize=(15, 10))
362 plt.subplots_adjust(hspace=0.4, wspace = 0.6)

```

```

363 for i, feature in enumerate(weather, 1):
364     plt.subplot(2, 4, i)
365     sns.countplot(x=feature, hue='Severity4', data=df_bl
        ↪ ,palette="Set2")
366
367     plt.xlabel('{}'.format(feature), size=12, labelpad=3)
368     plt.ylabel('Accident Count', size=12, labelpad=3)
369     plt.tick_params(axis='x', labelsiz=12)
370     plt.tick_params(axis='y', labelsiz=12)
371
372     plt.legend(['0', '1'], loc='upper right', prop={'size': 10})
373     plt.title('Count of Severity in \n {}
        ↪ Feature'.format(feature), size=14, y=1.05)
374 fig.suptitle('Count of Accidents by Weather Features (resampled
        ↪ data)', fontsize=18)
375 plt.show()
376
377 df = df.drop(['Heavy_Rain', 'Heavy_Snow', 'Fog'], axis = 1)
378
379 df = df.drop(['Wind_Direction'], axis=1)
380
381 POI_features =
        ↪ ['Amenity', 'Bump', 'Crossing', 'Give_Way', 'Junction', 'No_Exit', 'Railway', 'Ro
382
383 fig, axs = plt.subplots(ncols=3, nrow=4, figsize=(15, 10))
384
385 plt.subplots_adjust(hspace=0.5, wspace = 0.5)
386 for i, feature in enumerate(POI_features, 1):
387     plt.subplot(3, 4, i)
388     sns.countplot(x=feature, hue='Severity4', data=df_bl
        ↪ ,palette="Set2")
389
390     plt.xlabel('{}'.format(feature), size=12, labelpad=3)
391     plt.ylabel('Accident Count', size=12, labelpad=3)
392     plt.tick_params(axis='x', labelsiz=12)
393     plt.tick_params(axis='y', labelsiz=12)
394
395     plt.legend(['0', '1'], loc='upper right', prop={'size': 10})
396     plt.title('Count of Severity in {}'.format(feature), size=14,
        ↪ y=1.05)

```

```

397 fig.suptitle('Count of Accidents in POI Features (resampled
    ↳ data)',y=1.02, fontsize=16)
398 plt.show()
399
400 df =
    ↳ df.drop(['Amenity','Bump','Give_Way','No_Exit','Roundabout','Traffic_Calmi
    ↳ axis=1)
401
402 dtype_df = df_bl.dtypes.reset_index()
403 dtype_df.columns = ["Count", "Column Type"]
404 print(dtype_df)
405
406 # one-hot encoding
407 df[period_features] = df[period_features].astype('category')
408 df = pd.get_dummies(df, columns=period_features, drop_first=True)
409
410 # plot correlation
411 df_bl['Severity4'] = df_bl['Severity4'].astype(int)
412 plt.figure(figsize=(25,25))
413 cmap = sns.diverging_palette(220, 20, sep=20, as_cmap=True)
414 sns.heatmap(df_bl.corr(), annot=True,cmap=cmap,
    ↳ center=0).set_title("Correlation Heatmap", fontsize=14)
415 plt.show()
416
417 # Plot the proportion of severities.
418 plt.figure()
419 df_bl['Severity4'].value_counts().plot.pie(autopct='%1.1f%%')
420 plt.title('Percentage Severity Distribution')
421 plt.ylabel('Count')
422 plt.show()
423
424 df = df.drop(['Temperature(F)', 'Humidity(%)',
    ↳ 'Precipitation(in)', 'Precipitation_NA','Visibility_bc',
    ↳ 'Wind_Speed_bc',
425
    ↳ 'Clear','Cloud','Snow','Crossing','Junction','Railway','Mont
426 'Hour', 'Day','Minute',
    ↳ 'City_Freq','County_Freq','Airport_Code_Freq','Zipcode_Freq'
427 'Sunrise_Sunset_Night', 'Civil_Twilight_Night',
    ↳ 'Nautical_Twilight_Night'], axis=1)
428

```



```

429 df = df.drop(['Timezone'], axis=1)
430
431 # resample again
432 df_bl = resample(df, 'Severity4', 50000)
433
434 # plot correlation
435 df_bl['Severity4'] = df_bl['Severity4'].astype(int)
436 plt.figure(figsize=(20,20))
437 cmap = sns.diverging_palette(220, 20, sep=20, as_cmap=True)
438 sns.heatmap(df_bl.corr(), annot=True, cmap=cmap,
439             ↪ center=0).set_title("Correlation Heatmap", fontsize=14)
440 plt.show()
441
442 # Pre-process the dataset to extract the most important features
443 ↪ to predict the severity of an accident
444 # Find all continuous variables
445 continuous_vars = df.select_dtypes(include=['float64',
446 ↪ 'int64']).columns
447 print('The Dataset Contains, Continuous Variables:
448 ↪ {}'.format(continuous_vars))
449
450 # Find all categorical variables
451 categorical_vars = df.select_dtypes(include=['object', 'bool',
452 ↪ 'category']).columns
453 print('The Dataset Contains, Categorical Variables:
454 ↪ {}'.format(categorical_vars))
455
456 labels = []
457 values = []
458 for col in continuous_vars:
459     if col == 'Severity4':
460         continue
461     labels.append(col)
462     values.append(np.corrcoef(df[col].values,
463 ↪ df.Severity4.values)[0,1])
464 corr_df = pd.DataFrame({'col_labels':labels,
465 ↪ 'corr_values':values})
466 corr_df = corr_df.sort_values(by='corr_values')
467
468 ind = np.arange(len(labels))
469 width = 0.9

```

```

462 fig, ax = plt.subplots(figsize=(20,20))
463 rects = ax.barh(ind, np.array(corr_df.corr_values.values),
    ↪ color='b')
464 ax.set_yticks(ind)
465 ax.set_yticklabels(corr_df.col_labels.values,
    ↪ rotation='horizontal')
466 ax.set_xlabel("Correlation coefficient")
467 ax.set_title("Correlation coefficient of the variables")
468 plt.show()
469
470 from sklearn.model_selection import train_test_split
471 from sklearn.preprocessing import StandardScaler
472 from sklearn.decomposition import PCA
473 from sklearn.linear_model import LogisticRegression # You can
    ↪ use other models depending on your problem.
474 from sklearn.metrics import classification_report,
    ↪ confusion_matrix
475 from sklearn.metrics import confusion_matrix,
    ↪ ConfusionMatrixDisplay
476 from sklearn.impute import SimpleImputer
477 from sklearn.ensemble import AdaBoostClassifier,
    ↪ RandomForestClassifier
478 from sklearn.model_selection import GridSearchCV, KFold,
    ↪ train_test_split, cross_val_predict
479 from sklearn.metrics import classification_report,
    ↪ confusion_matrix
480 x_train, x_test, y_train, y_test =
    ↪ train_test_split(df_bl.drop('Severity4', axis=1),
    ↪ df_bl['Severity4'], test_size=0.2, random_state=42)
481
482 my_imputer = SimpleImputer()
483 x_train = my_imputer.fit_transform(x_train)
484
485 # Perform PCA analysis with Skree plot to find the most important
    ↪ features
486 # Standardize the data
487 scaler = StandardScaler()
488 df_scaled = scaler.fit_transform(x_train)
489 print(df_scaled)
490
491 print(x_train.shape)

```

```

492 # Perform PCA
493 pca = PCA(n_components=8)
494 df_pca = pca.fit_transform(df_scaled)
495
496 # Plot the Skree plot
497 plt.figure(figsize=(12, 10))
498 plt.plot(np.cumsum(pca.explained_variance_ratio_), 'bo-')
499 plt.grid()
500 plt.xlabel('Number of Components')
501 plt.ylabel('Cumulative Explained Variance Ratio')
502 plt.title('Scree Plot')
503 plt.show()
504
505 x_test = my_imputer.fit_transform(x_test)
506 df_scaled_test = scaler.transform(x_test)
507 df_pca_test = pca.transform(df_scaled_test)
508 print(df_pca_test)
509
510 # Initialize and train a machine learning model (e.g., Logistic
    ↳ Regression)
511 model = LogisticRegression(max_iter=1000)
512 model.fit(x_train, y_train)
513
514 # Make predictions on the test set
515 y_pred = model.predict(x_test)
516
517 # Calculate the accuracy of the model
518 print(classification_report(y_test, y_pred))
519
520 confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)
521
522 conf_matrix = pd.DataFrame(data=confmat,
523
    ↳ columns=['Predicted:0', 'Predicted:1'], index=['A
524 plt.figure(figsize = (8,5))
525 sns.heatmap(conf_matrix,
    ↳ annot=True,fmt='d',cmap="YlGnBu").set_title(
526     "Confusion Matrix \n Logistic Regression", fontsize=16)
527 plt.show()
528
529 clf_base = RandomForestClassifier()

```

```

530 grid = {'n_estimators': [10, 50, 100],
531         'max_features': ['auto', 'sqrt']}
532 clf_rf = GridSearchCV(clf_base, grid, cv=5, n_jobs=8,
533                       ↪ scoring='f1_macro')
534
535 clf_rf.fit(x_train, y_train)
536 y_pred = clf_rf.predict(x_test)
537
538 print(classification_report(y_test, y_pred))
539
540 confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)
541
542 conf_matrix = pd.DataFrame(data=confmat,
543                             ↪ columns=['Predicted:0', 'Predicted:1'], index=['A
544 plt.figure(figsize = (8,5))
545 sns.heatmap(conf_matrix,
546             ↪ annot=True,fmt='d',cmap="YlGnBu").set_title(
547         "Confusion Matrix \n Random Forest", fontsize=16)
548 plt.show()
549
550 importances = pd.DataFrame(np.zeros((x_train.shape[1], 1)),
551                             ↪ columns=['importance'],
552                             ↪ index=df.drop('Severity4',axis=1).columns)
553
554 importances.iloc[:,0] =
555     ↪ clf_rf.best_estimator_.feature_importances_
556
557 importances.sort_values(by='importance', inplace=True,
558     ↪ ascending=False)
559 importances30 = importances.head(30)
560
561 plt.figure(figsize=(15, 10))
562 sns.barplot(x='importance', y=importances30.index,
563     ↪ data=importances30)
564
565 plt.xlabel('')
566 plt.tick_params(axis='x', labels=10)
567 plt.tick_params(axis='y', labels=10)
568 plt.title('Random Forest Classifier Feature Importance', size=15)
569

```

```
563 | plt.show()
```